

Database Storage Engines (Key-Value Stores or Hash Tables)

Hung-Chang Hsiao (蕭宏章)

Professor

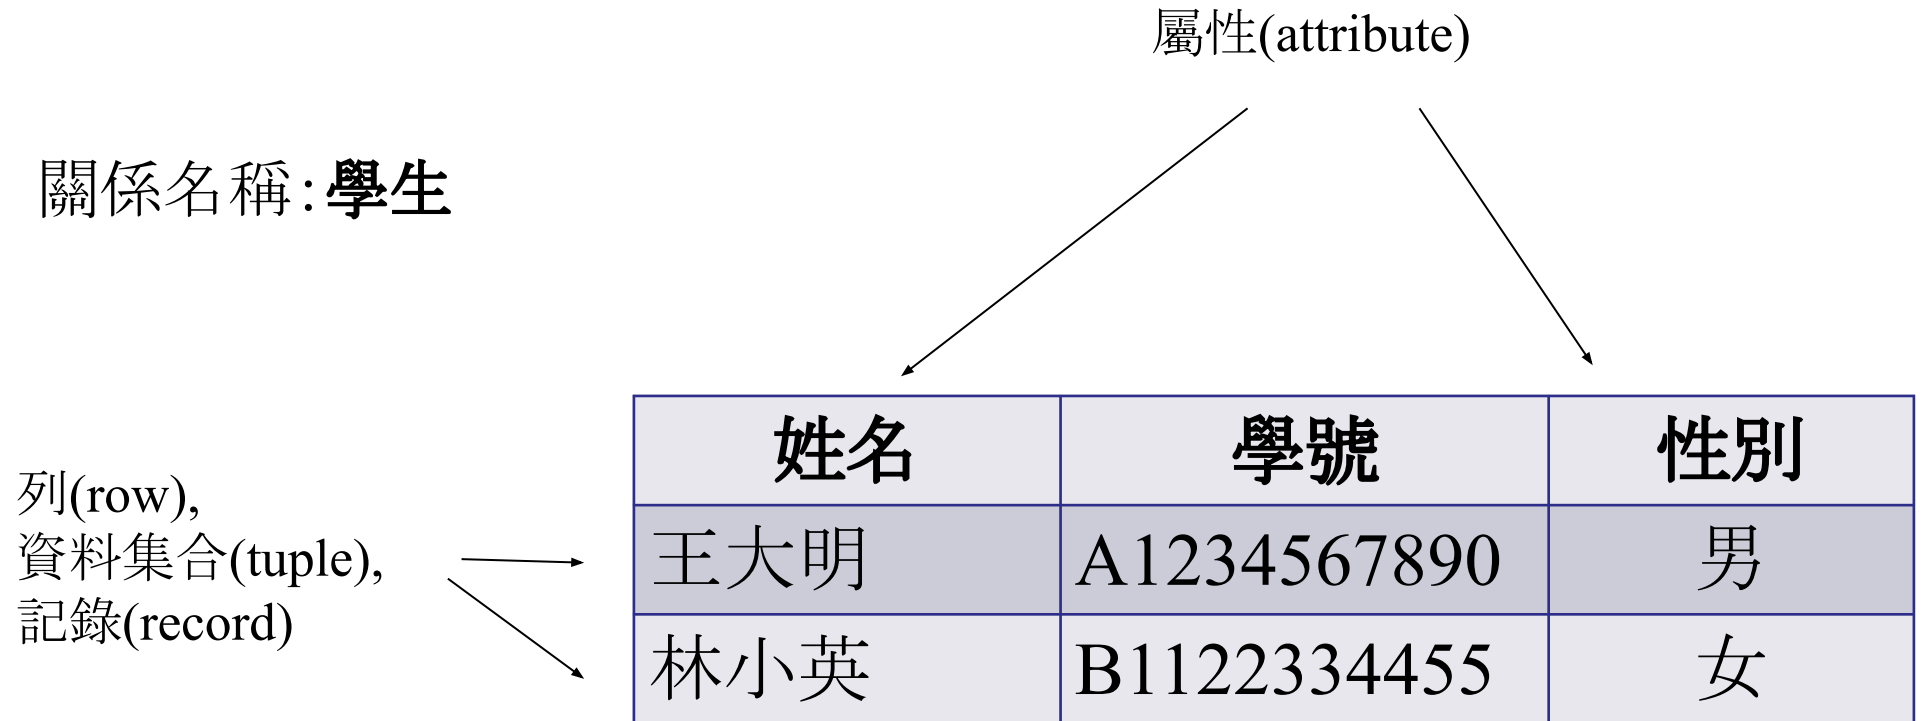
Department of Computer Science and Information Engineering
National Cheng Kung University
Taiwan

Databases in a Word

- **SQL compiler + Database storage engine**
 - SQL compiler: 將高階的SQL語言換成低階的storage engine語言
 - e.g., "SELECT * from Table XYZ" -> "SCAN (start key, end key)"
 - Database storage engine: 執行低階的資料操作語言並與檔案系統互動
- **MariaDB (i.e., MySQL) 在給定的SQL compiler下得置換 storage engine**

如何表示關係？

■ 關係是由資料表所構成



資料庫綱要 (Database Schema)

- 綱要(Schema)

- 關係名稱和屬性所組成
- 描述資料的形態
- E.g., Student(name, studentID) 或是 Student(name: string, studentID: string)

資料庫定義語言

(Data-Definition Language, or DDL)

- 用以描述, 修改或是刪除資料庫中的資料關係
- 定義欄位、資料型態、資料結構
- E.g.,

```
CREATE TABLE STUDENT(  
    name CHAR(10),  
    studentID CHAR(15),  
    gender CHAR(1),  
    PRIMARY KEY (studentID)  
)
```

資料庫處理語言

(Data-Management Language, or DML)

- **讓使用者存取或是處理資料庫的資料**

- 讀取資料
- 新增資料
- 修改更新資料
- 刪除資料

- **E.g.,**

```
INSERT INTO STUDENT(name, studentID, gender)  
VALUES('David', 'A1231231234', 'M')
```

NoSQL Databases (Storage Engines)

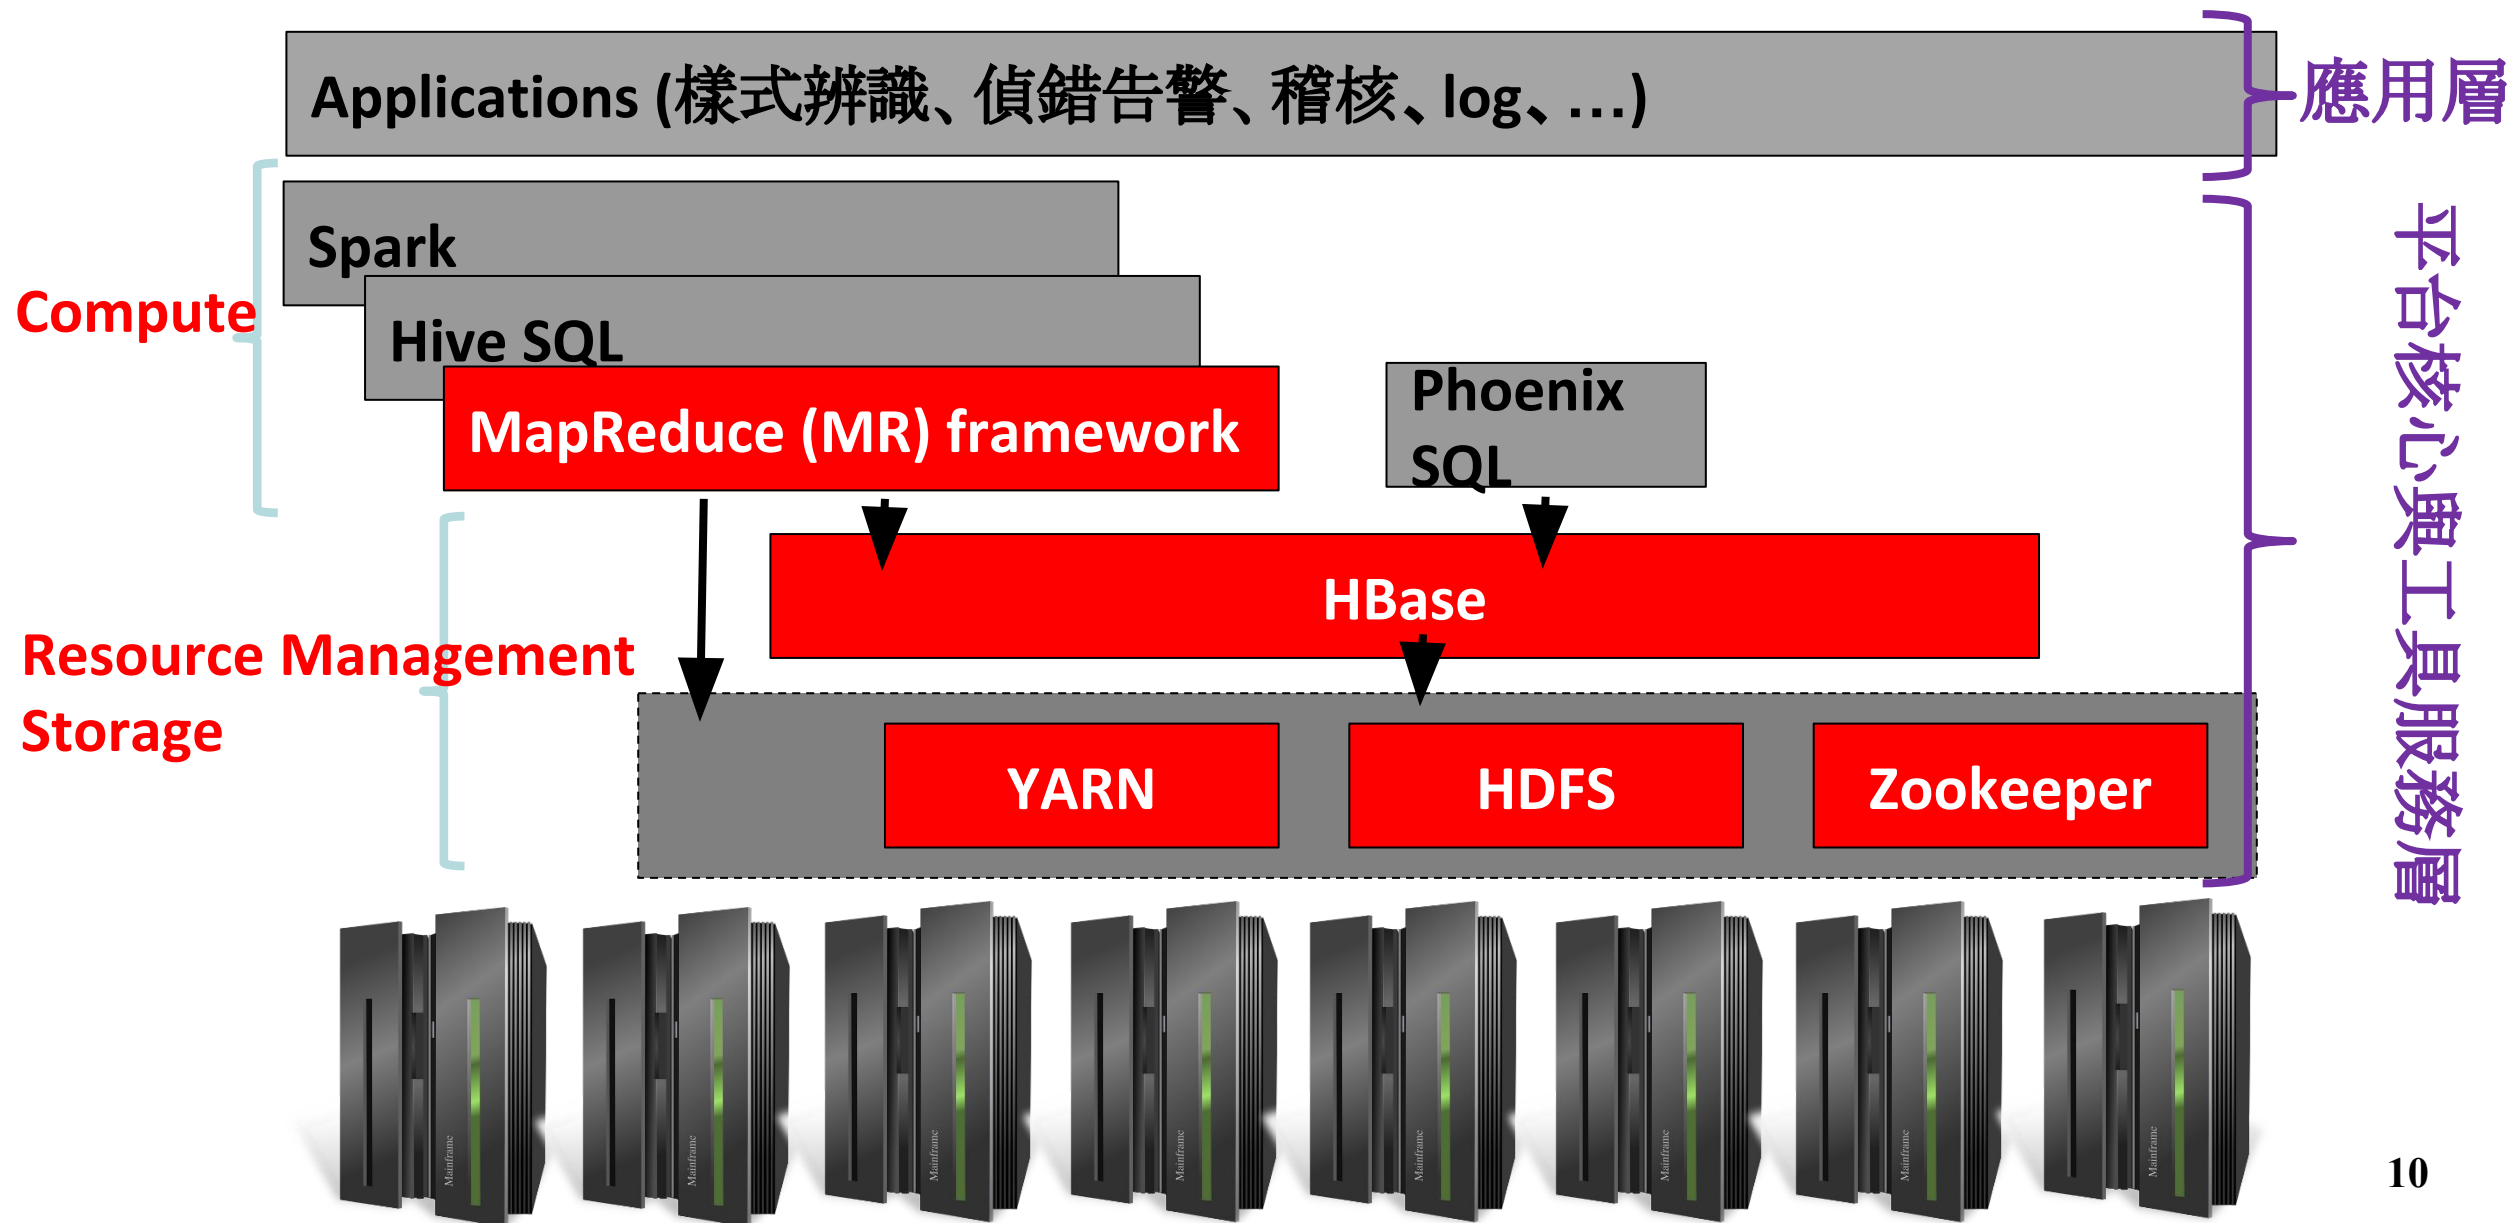
- No relations among DB tables
 - DB relations across tables 是透過 foreign keys 來相互關聯表格的屬性欄位關係
 - Relations 正是讓 DB engine 無法橫向擴充 (scale-out) 的關鍵原因
- APIs
 - GET (key)
 - PUT (key, value)
 - SCAN (start key, end key)
- Highly scalable: distributed DB engines
- SQL-like support, e.g., Apache Hive and Phoenix over HBase
- No transactions
- 經驗: 無論 SQL 或 NoSQL, 當對 DB 執行速度有嚴謹的要求時, 我們會直接操作 storage engine 而非透過 SQL 語言

Question?

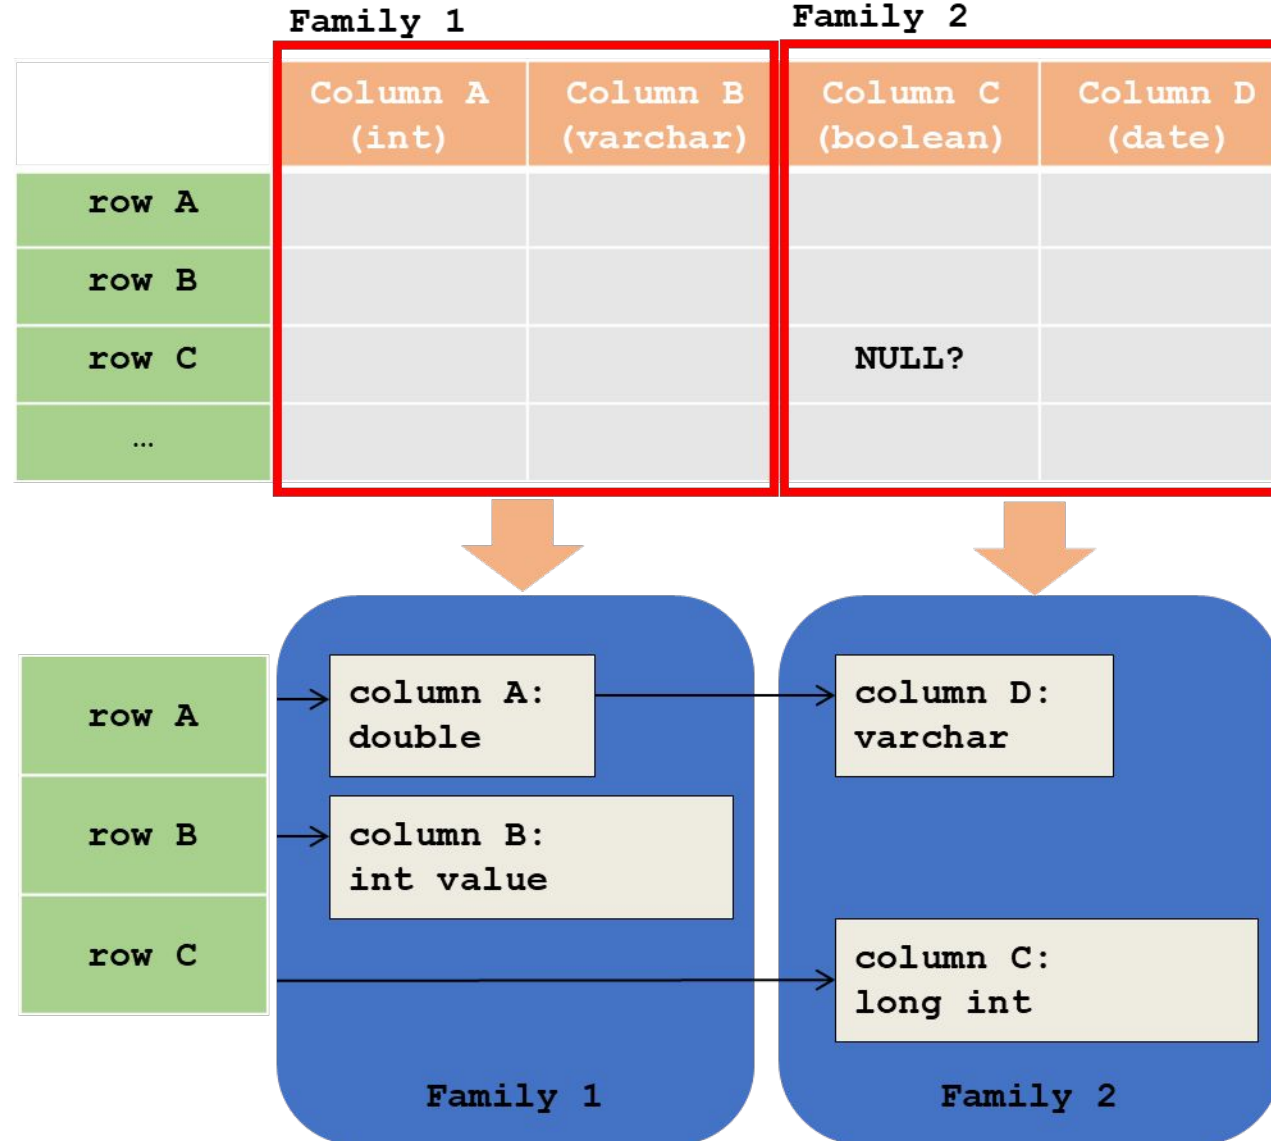
- 給定大量key-value pairs的資料, 我們想查找給定符合某些條件keys的values, 請問如何做?
 - E.g., (1, a), (2, b), (3, c), ..., (10, j)
 - 找出 $2 < \text{keys} < 4$ 的values值
- 要求: 我們不能控制key-value pairs資料產生的keys值順序!
- 如何儲存該些資料? 若只靠檔案系統 (無論集中或分散)?
 - 假設資料來的順序是 (3, c), (10, j), (2, b), (1, a), ...
 - 我們按資料產生的順序來將資料儲存在數個檔案裡, 比如檔案A存 (3, c) 及 (10, j)、檔案B存 (2, b) 及 (1, a), 檔案C存...
- 如何儲存會大大影響將來的資料查找的效率



Big Data Platform: Hadoop Ecosystem



Overview: HBase Table



API Overview (1/5): Put Data

Class **Configurations**

- Configurations are specified by resources.
- A resource contains a set of name/value pairs as XML data

```
Configuration conf = HBaseConfiguration.create();
try (Connection connection = ConnectionFactory.createConnection(conf)) {
    try (Table table = connection.getTable(TableName.valueOf("table name"))) {
        List<Put> putList = new LinkedList();
        Put put1 = new Put(Bytes.toBytes("row1"));
        put1.addColumn(Bytes.toBytes("column"),
            Bytes.toBytes("qualifier"),
            Bytes.toBytes("value"));
        Put put2 = new Put(Bytes.toBytes("row2"));
        Cell cell = CellUtil.createCell(
            Bytes.toBytes("column"),
            Bytes.toBytes("column"),
            Bytes.toBytes("column"),
            System.currentTimeMillis(),
            KeyValue.Type.Put.getCode(),
            Bytes.toBytes("column"));
        put2.add(cell);
        putList.add(put1);
        putList.add(put2);
        table.put(putList);
    }
}
```

Interface **Connection**

- A cluster connection encapsulating lower level individual connections to actual servers and a connection to zookeeper.

Interface **Table**

- Used to communicate with a single HBase table

Class **Put**

- Used to perform Put operations for a single row
- A put is composed of many cells

Interface **Cell**

- The unit of storage in HBase consisting of the following fields

1. **Row**
2. **column family**
3. **column qualifier**
4. **Timestamp**
5. **Type**
6. **MVCC version (set by server)**
7. **value**

API Overview (2/5): Delete Data

```
Configuration conf = HBaseConfiguration.create();
try (Connection connection = ConnectionFactory.createConnection(conf)) {
    try (Table table = connection.getTable(TableName.valueOf("table name"))) {
        List<Delete> deleteList = new LinkedList();
        Delete delete1 = new Delete(Bytes.toBytes("row1"));
        delete1.addColumn(Bytes.toBytes("column"),
            Bytes.toBytes("qualifier"));
        Delete delete2 = new Delete(Bytes.toBytes("row2"));
        Cell cell = CellUtil.createCell(
            Bytes.toBytes("column"),
            Bytes.toBytes("column"),
            Bytes.toBytes("column"),
            System.currentTimeMillis(),
            KeyValue.Type.Delete.getCode(),
            Bytes.toBytes("column"));
        delete2.addDeleteMarker(cell);
        deleteList.add(delete1);
        deleteList.add(delete2);
        table.delete(deleteList);
    }
}
```

Class **Delete**

- To delete an entire row, instantiate a Delete object with the row to delete.
- To define the scope of what to delete

API Overview (3/5): Get Data

```
Configuration conf = HBaseConfiguration.create();
try (Connection connection = ConnectionFactory.createConnection(conf)) {
    try (Table table = connection.getTable(TableName.valueOf("table name"))) {
        Get get = new Get(Bytes.toBytes("row1"));
        Result rowResult = table.get(get);
        for (Cell cell : rowResult.rawCells()) {
        }
        Scan scan = new Scan();
        try (ResultScanner scanner = table.getScanner(scan)) {
            for (Result result : scanner) {
                for (Cell cell : result.rawCells()) {
                }
            }
        }
    }
}
```

Class **Get**

- Used to perform Get operations on a single row.

Class **Result**

- Single row result of a query
- A result is composed of many cells

Class **Scan**

- Used to perform Scan operations.
- Rather than specifying a single row, an optional **startRow** and **stopRow** may be defined

Interface **ResultScanner**

- iterate over all rows.

API Overview (4/5): Create a Table

Interface **Admin**

- The administrative API for HBase

```
Configuration conf = HBaseConfiguration.create();
try (Connection connection = ConnectionFactory.createConnection(conf)) {
    try (Admin admin = connection.getAdmin()) {
        HTableDescriptor tableDesc
            = new HTableDescriptor(TableName.valueOf("table name"));
        HColumnDescriptor columnDesc
            = new HColumnDescriptor(Bytes.toBytes("column"));
        tableDesc.addFamily(columnDesc);
        admin.createTable(tableDesc);
    }
}
```

class **HColumnDescriptor**

- contains information about a column family
1. Set block cache
 2. Set bloom filter
 3. Compression
 4. Data block encoding
 5. HDFS block size
 6. Max/min version

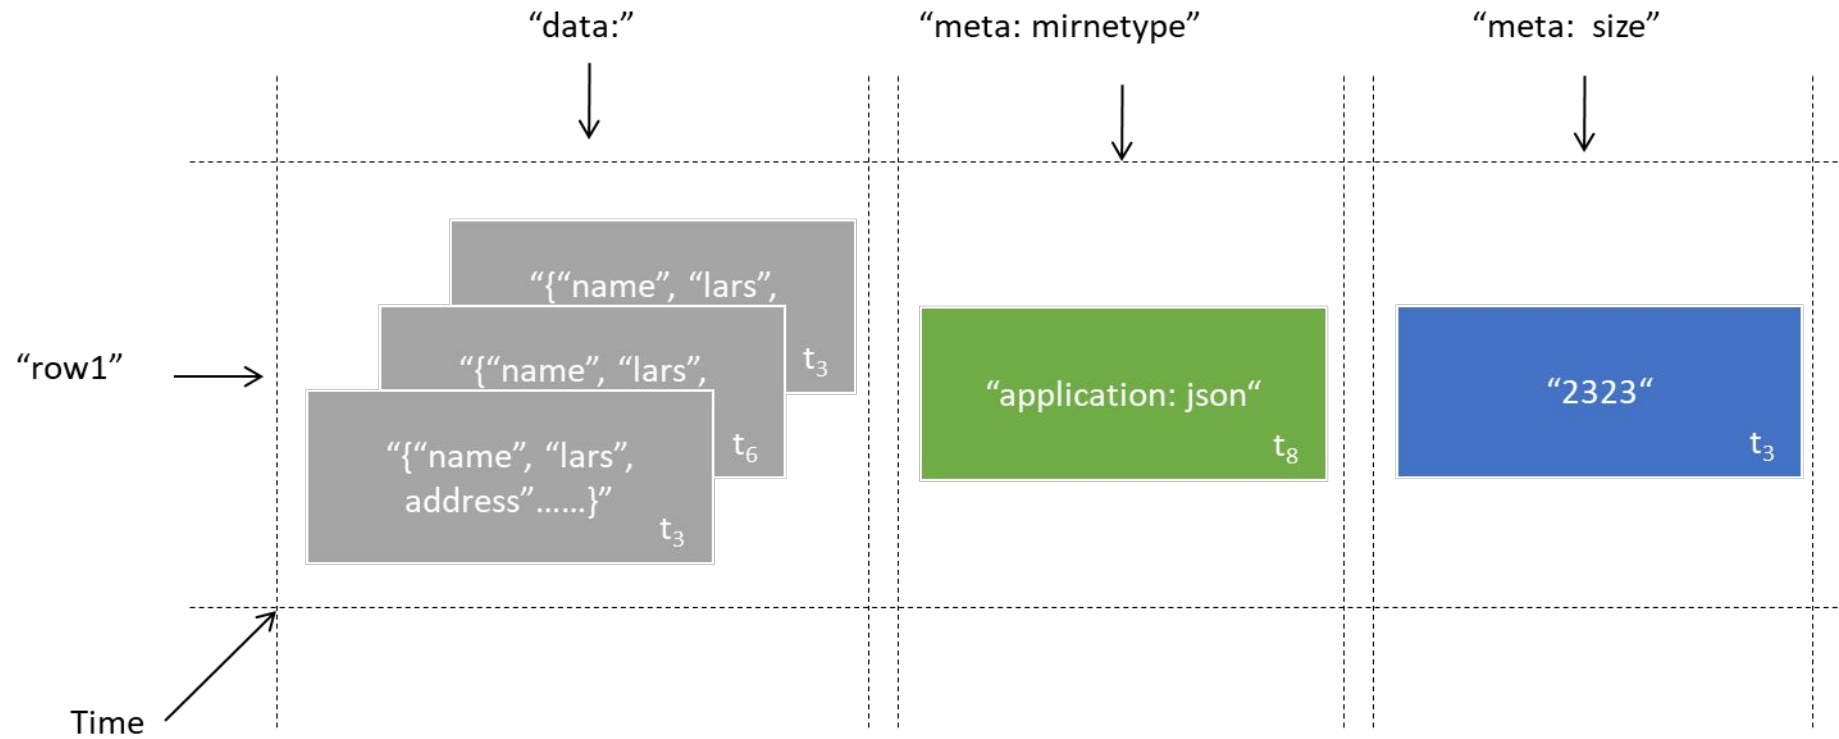
class **HTableDescriptor**

- contains the details about an HBase table
1. Column
 2. Compaction
 3. Durability
 4. File size
 5. Memstore size

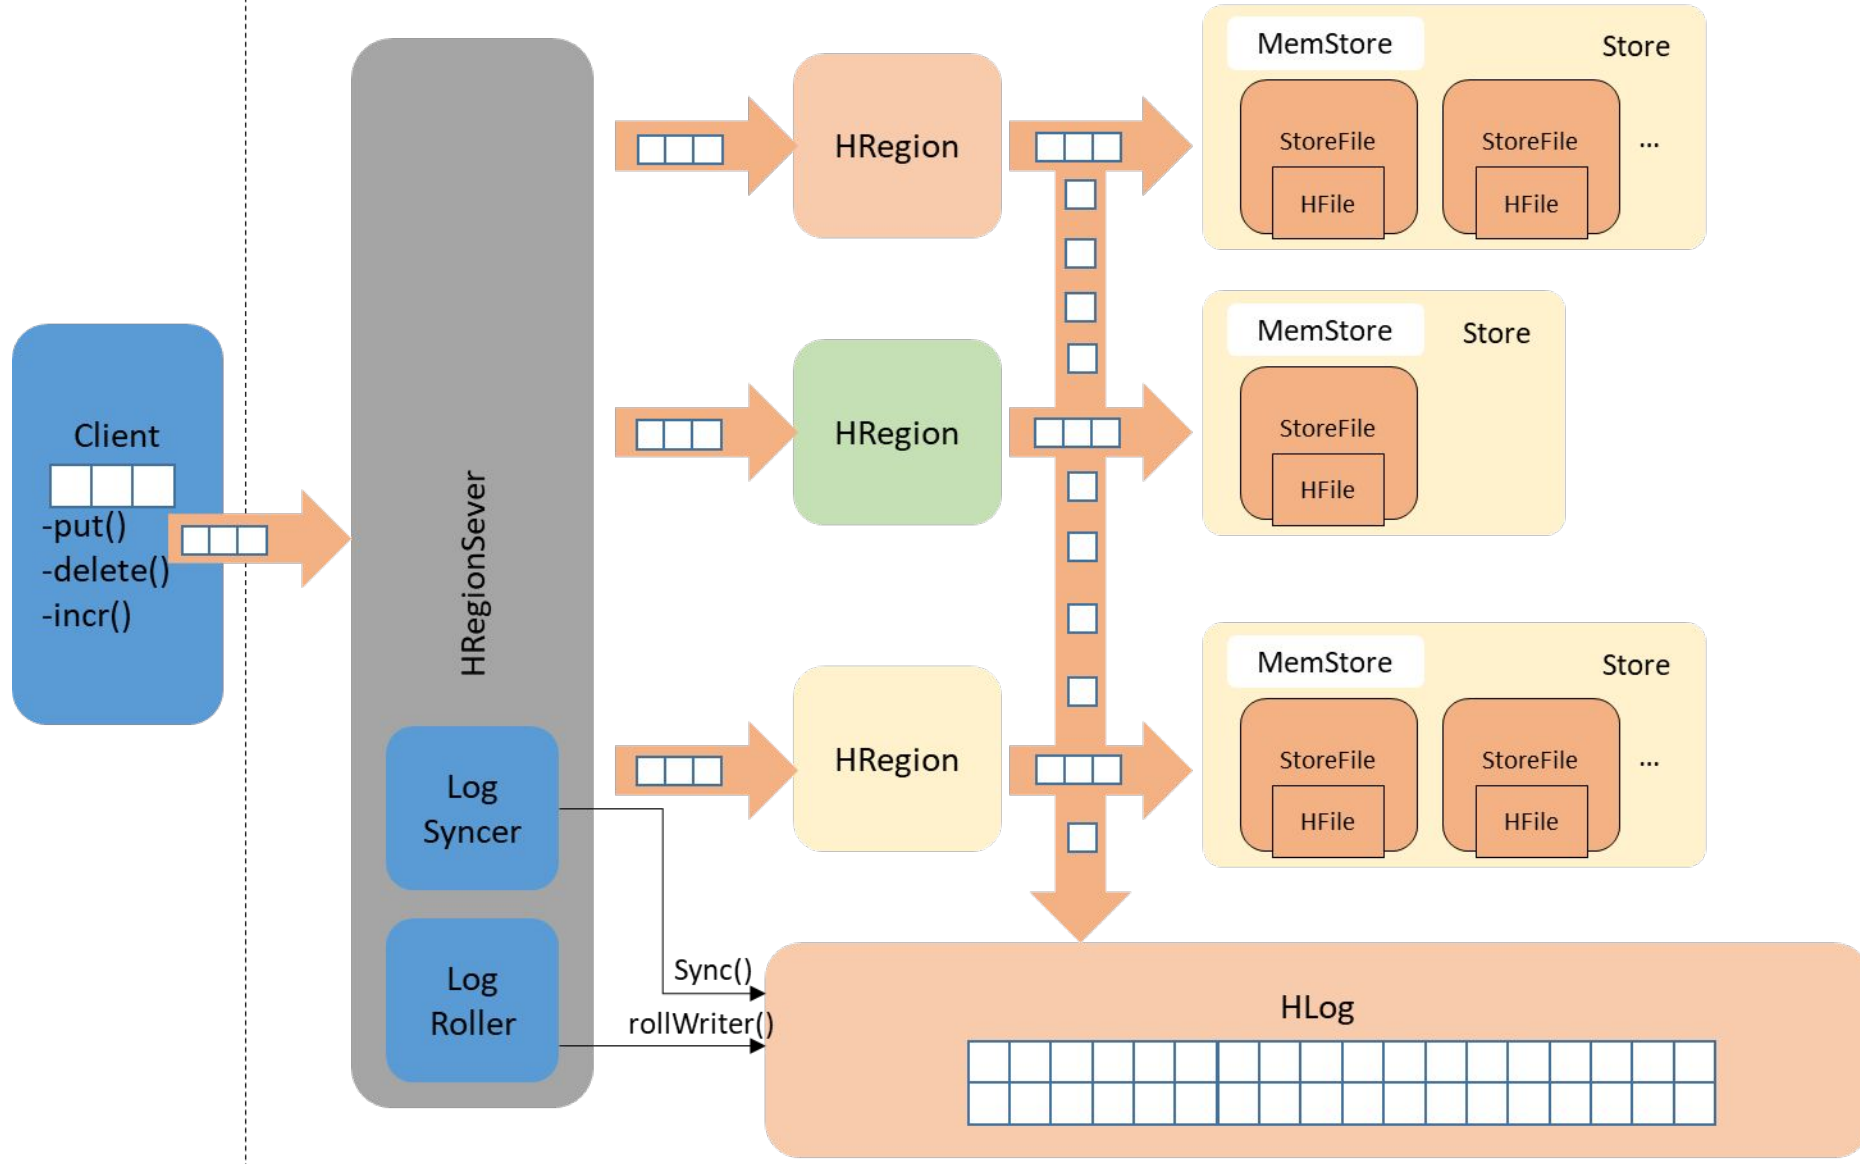
API Overview (5/5): Admin Interface

Method Name	Description
<code>void modifyColumn(final TableName tableName, final HColumnDescriptor descriptor) throws IOException;</code>	Modify a column
<code>void modifyTable(final TableName tableName, final HTableDescriptor htd) throws IOException;</code>	Modify a table
<code>void deleteColumn(final TableName tableName, final byte[] columnName) throws IOException;</code>	Deletes a column
<code>void deleteTable(final TableName tableName) throws IOException;</code>	Deletes a table
<code>void flush(final TableName tableName) throws IOException;</code>	Flush a table
<code>ClusterStatus getClusterStatus() throws IOException;</code>	Return the cluster status
<code>void enableTable(final TableName tableName) throws IOException;</code>	Enable table and wait on completion
<code>void disableTable(final TableName tableName) throws IOException;</code>	Disable table and wait on completion
<code>HTableDescriptor[] listTables() throws IOException;</code>	List all the userspace tables.
<code>void majorCompact(TableName tableName) throws IOException;</code>	Major compact a table
<code>void mergeRegions(final byte[] encodedNameOfRegionA, final byte[] encodedNameOfRegionB, final boolean forcible) throws IOException;</code>	Merge two regions

Timestamp



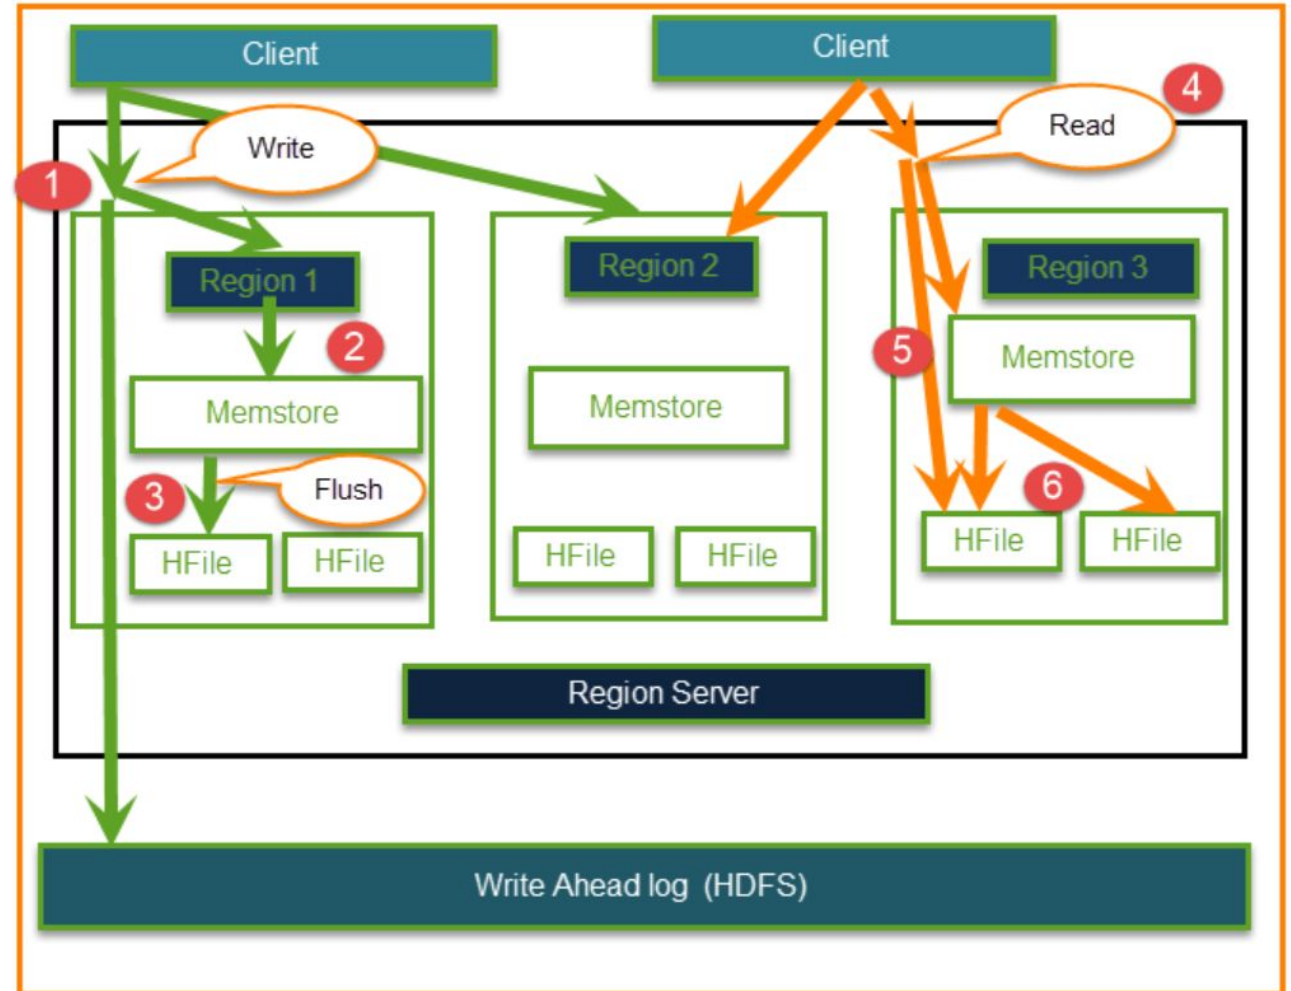
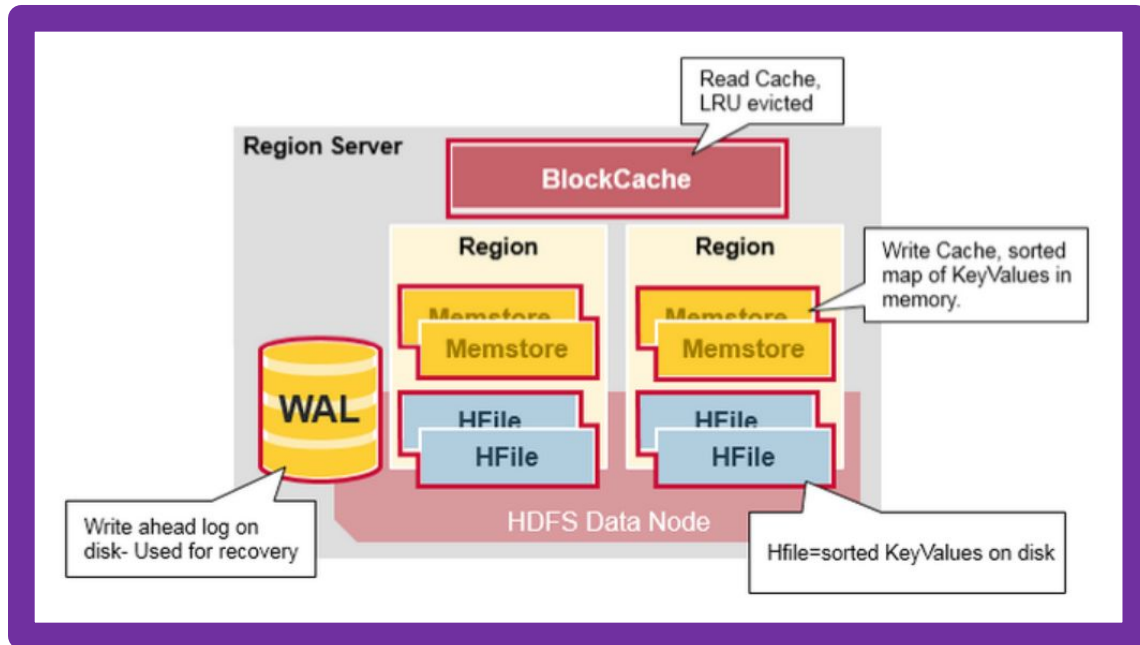
Architecture: MemStore, HLog and HFile



Write-Ahead Log (HLog)

- Region servers keep data in-memory until enough is collected to warrant a flush to disk, avoiding the creation of too many very small files
- By default, each in-memory update is written to a log
- Available options:
 - **ASYNC_WAL**: Write the Mutation to the WAL asynchronously
 - **FSYNC_WAL**: Write the Mutation to the WAL synchronously and force the entries to disk
 - **SKIP_WAL**: Do not write the Mutation to the WAL
 - **SYNC_WAL** (default): Write the Mutation to the WAL synchronously

Get: Data Flow

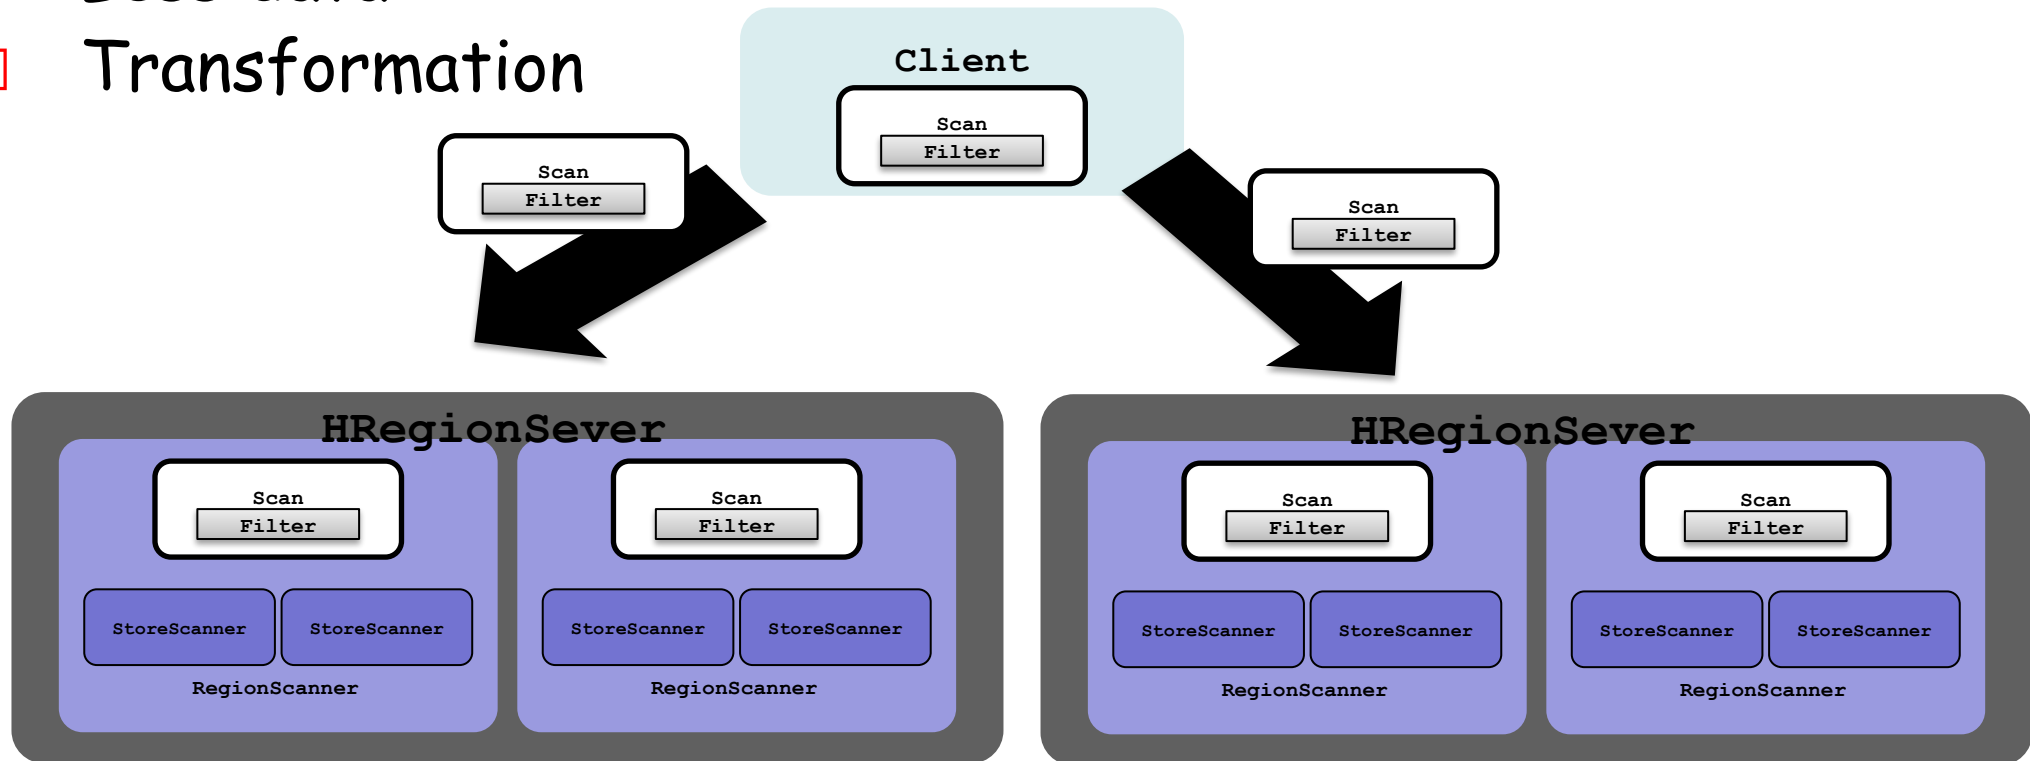


Splitting a Region

- A split policy determines how a region should be split
 - The fundamental unit of split is "row"
- Various options:
 - configured in `hbase.regionserver.region.split.policy`
 - `ConstantSizeRegionSplitPolicy` class
 - `IncreasingToUpperBoundRegionSplitPolicy` class (default)
 - `DelimitedKeyPrefixRegionSplitPolicy` class
 - `KeyPrefixRegionSplitPolicy` class
 - `DisabledRegionSplitPolicy` class

Scanner and Filter

- Used in GET () and SCAN () API calls
- Fine-grained control over what is returned to the client
 - Less data
 - Transformation



Available Filters (cont'd)

Class Name	Description
<code>ColumnCountGetFilter</code>	<ul style="list-style-type: none">• Return first <code>N</code> columns on row only
<code>ColumnPaginationFilter</code>	<ul style="list-style-type: none">• Based on the <code>ColumnCountGetFilter</code>, takes two arguments: limit and offset
<code>ColumnPrefixFilter</code>	<ul style="list-style-type: none">• Select only those keys with columns that match a particular prefix
<code>MultipleColumnPrefixFilter</code>	<ul style="list-style-type: none">• Select only those keys with columns that match a set of prefixes
<code>ColumnRangeFilter</code>	<ul style="list-style-type: none">• Select only those keys with columns that are between <code>minColumn</code> to <code>maxColumn</code>• Can specify an operator (equal, greater, not equal, etc)
<code>CompareFilter</code>	<ul style="list-style-type: none">• To filter by row key, use <code>RowFilter</code>.• To filter by column qualifier, use <code>QualifierFilter</code>.• To filter by value, use <code>SingleColumnValueFilter</code>.
<code>FirstKeyOnlyFilter</code>	<ul style="list-style-type: none">• Return first KV from each row
<code>MultiRowRangeFilter</code>	<ul style="list-style-type: none">• Scan multiple row key ranges

Available Filters (cont'd)

Class Name	Description
<code>FuzzyRowFilter</code>	<ul style="list-style-type: none">Specify (row key, fuzzy info) to match row keys, where fuzzy info equal to<ul style="list-style-type: none">0: not interested in a particular byte in a key1: interested in a particular byte in a key
<code>InclusiveStopFilter</code>	<ul style="list-style-type: none">Stop once touching a given row
<code>KeyOnlyFilter</code>	<ul style="list-style-type: none">Return key of each KV
<code>PageFilter</code>	<ul style="list-style-type: none">Limit results in a specific page size
<code>PrefixFilter</code>	<ul style="list-style-type: none">Result values that have same row prefix
<code>RandomRowFilter</code>	<ul style="list-style-type: none">Rows that are interested with a probability
<code>SingleColumnValueFilter</code>	<ul style="list-style-type: none">Filter cells based on value
<code>SkipFilter</code>	<ul style="list-style-type: none">Filter an entire row if any of the Cell checks fails

Filter on Your Own

Method Name	Description
<code>abstract public void reset() throws IOException;</code>	Reset the state of the filter between rows
<code>abstract public boolean filterRowKey(byte[] buffer, int offset, int length) throws IOException;</code>	Filters a row based on row key
<code>abstract public boolean filterAllRemaining() throws IOException;</code>	If true, the scan will terminate
<code>abstract public ReturnCode filterKeyValue(final Cell v) throws IOException;</code>	Way to filter based on the column family, column qualifier and/or the column value
<code>abstract public Cell transformCell(final Cell v) throws IOException;</code>	Give the filter a chance to transform the passed Key-Value
<code>abstract public void filterRowCells(List<Cell> kvs) throws IOException;</code>	Alter the contains of specified Cells
<code>abstract public boolean hasFilterRow();</code>	Primarily used to check for conflicts with scans such as scans that do not read a full row at a time
<code>abstract public boolean filterRow() throws IOException;</code>	Last chance to filter row based on previous <code>filterKeyValue(Cell)</code> calls
<code>abstract public Cell getNextCellHint(final Cell currentCell) throws IOException;</code>	If the filter returns <code>SEEK_NEXT_USING_HINT</code> , then it should also tell which is the next key it must seek to
<code>abstract public boolean isFamilyEssential(byte[] name) throws IOException;</code>	Check a given column family whether is essential to filter
<code>abstract public byte[] toByteArray() throws IOException;</code>	Serialize filter to byte array
<code>public static Filter parseFrom(final byte [] pbBytes) throws DeserializationException</code>	Failure signal

Coprocessor

- An alike MapReduce framework that distributes work across the entire cluster (such as filters)
 - Enable to run arbitrary code directly on each region server
- Types of coprocessor:

Observer: comparable to triggers (or callback functions) that are executed when certain events occur

- **RegionObserver** class
- **MasterObserver** class
- **WALObserver** class

Endpoint: user code can be deployed to the servers hosting the data to perform server-local computations

More Features

Feature	Description
Replication	<ul style="list-style-type: none">• copy data between HBase deployments
Bloom filter	<ul style="list-style-type: none">• predict whether a given element is a member in a set of data
Metrics	<ul style="list-style-type: none">• expose a large number of metrics that detail present status
Compression	<ul style="list-style-type: none">• compression algorithms for an Hfile
Compaction	<ul style="list-style-type: none">• combine HFiles to a few, larger Hfiles
MapReduce Integration	TableInputFormat <ul style="list-style-type: none">• Convert HBase tabular data into a format that can be understood by Map/Reduce
	TableOutputFormat <ul style="list-style-type: none">• Convert Map/Reduce output to an HBase table

Compactions (Per Region Based)

- **Minor compaction**

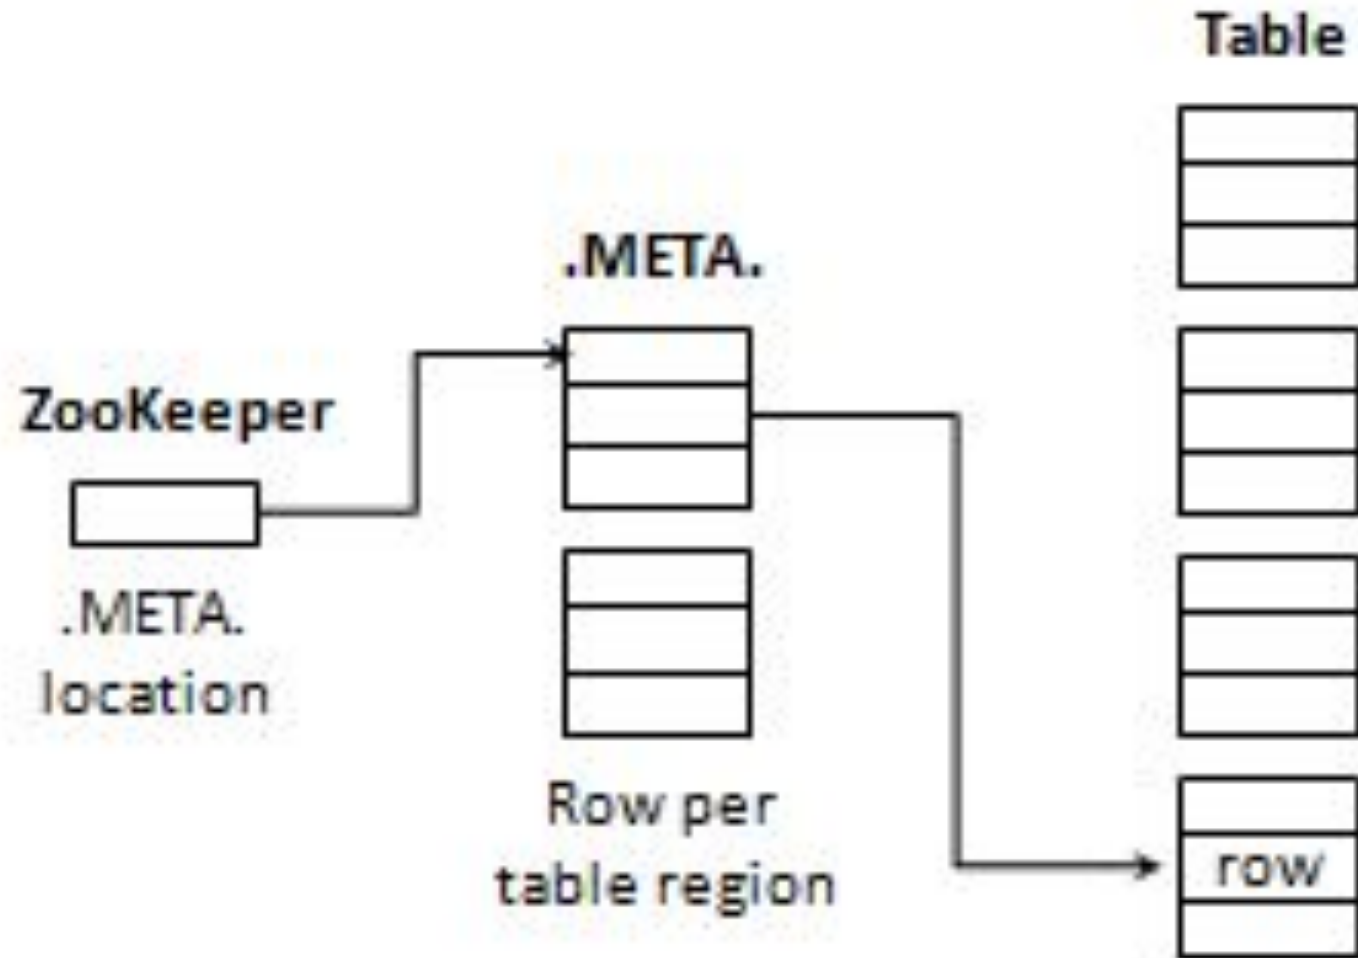
- 將幾個小的Hfiles整理成一個大的Hfile (to minimize the seek time)

- **Major compaction**


- HBase特別適合大量資料的寫入，持續產生Hfiles，致使反覆被update的資料衍生數個版本分佈在不同的HFiles裡；同時也移除delete marks!
- 浪費儲存空間，也增加搜尋資料的時間
- 希望“所有HFiles”裡每筆相異資料的版本個數為使用者所指定，同時也希望所有的Hfiles裡的已經按照keys的global order整理

- **Compaction is a background process**

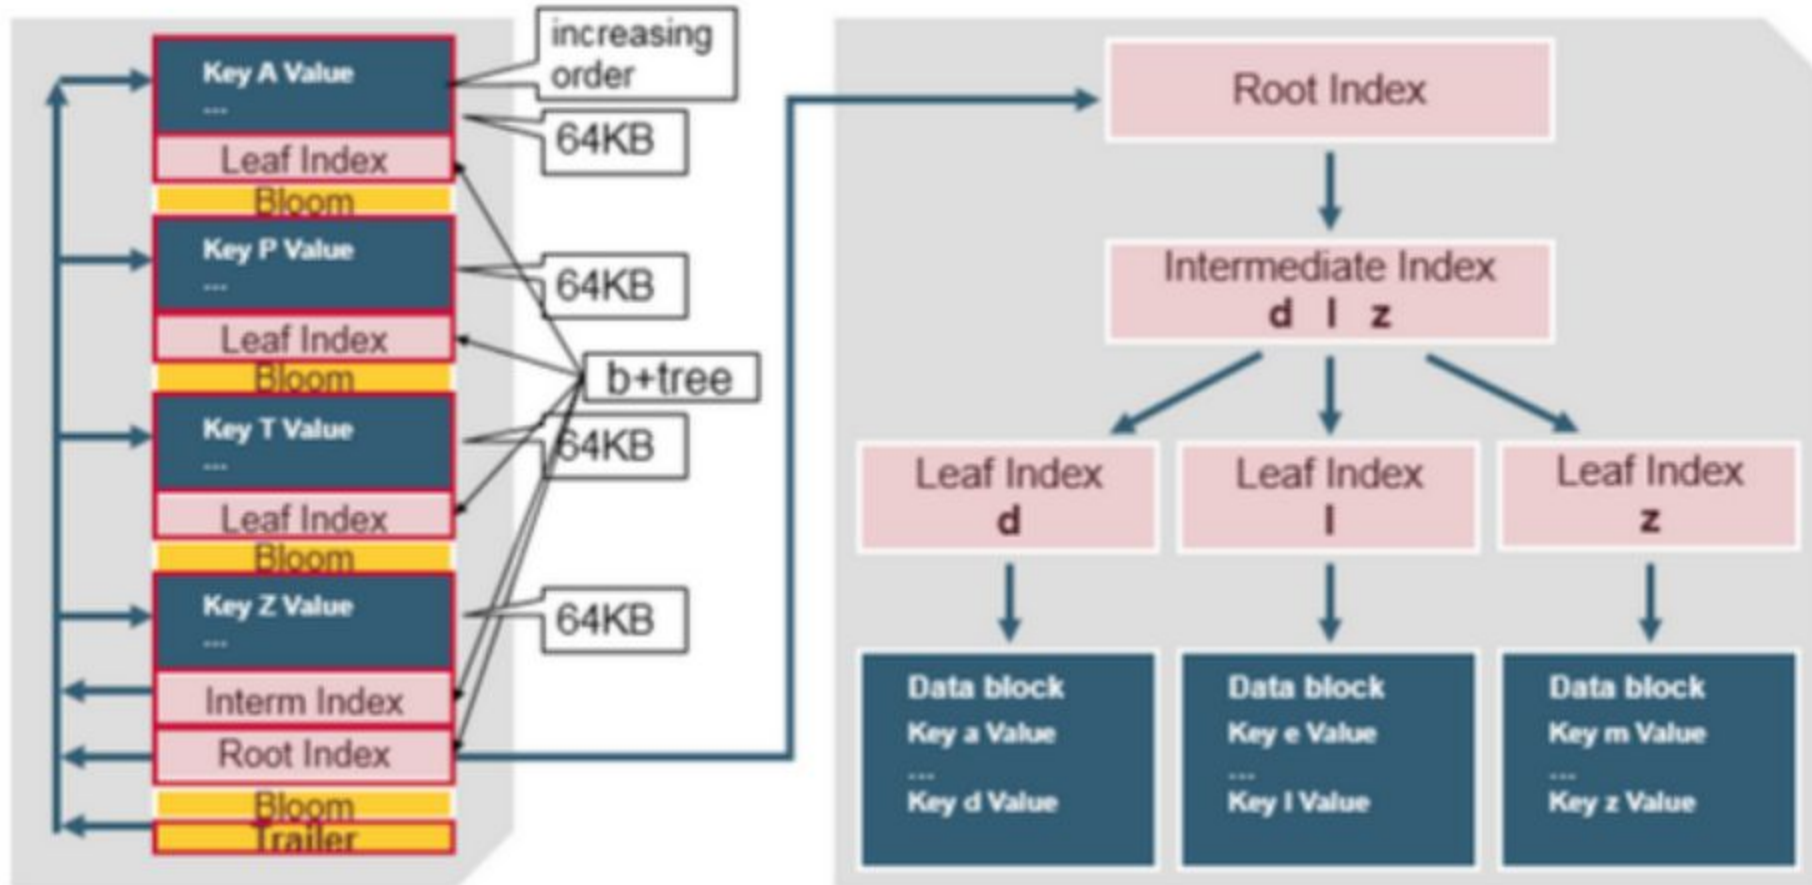
.META Table



.META Table (cont'd)



HFile Structure



Bloom Filters

- A concise representation of a set S with m elements using a n -bit vector, given:
 - Denote by $S = \{x_1, x_2, x_3, \dots, x_m\}$
 - k hash functions (denoted by H_i , where $1 \leq i \leq k$), each picking a value from $\{1, 2, 3, \dots, n\}$ uniformly at random
- The n -bit **standard Bloom filter** (\mathbb{B}) is initially set to $0_1 0_2 0_3 \dots 0_n$
- Insert an element y into \mathbb{B} : set bit $H_i(y) = 1$ for all $i = 1, 2, \dots, k$

Example

$1_1 0_2 0_3 1_4 1_5$ for $n = 5$ and $k = 3$

- An element $y \in \mathbb{B}$ if $H_i(y) = 1$ for all $i = 1, 2, \dots, k$

Bloom Filters (cont'd)

- **False positive:** An element $y \notin \mathbb{B}$ is claimed in \mathbb{B}
- The probability of a specific bit is zero is

$$p = \left(1 - \frac{1}{n}\right)^{km} \approx e^{-\frac{km}{n}} \quad (1)$$

- The false positive rate f is

$$f = (1 - p)^k \quad (2)$$

Example

$f \approx 0.02$ if $n = 8m$ and $k = 5$ (or $k = 6$)

Bloom Filters (cont'd)

- Observation:

- Larger k , more chance to find a 0-bit for $y \notin S$
- Smaller k , more fraction of 0 bit and thus smaller f

Theorem

Given n and m , k minimizes f if

$$k = \frac{n \ln 2}{m}. \quad (3)$$

Proof.

- 1 $f = (1 - p)^k = e^g$, where $g = k \ln(1 - p)$.
- 2 Minimizing f is identical to minimize g , and let $\frac{dg}{dk} = 0$.



Bloom Filters (cont'd)

Corollary

If $k = \frac{n \ln 2}{m}$, then $p = \frac{1}{2}$.

■ That is,

$$f = \left(\frac{1}{2}\right)^k \quad (4)$$

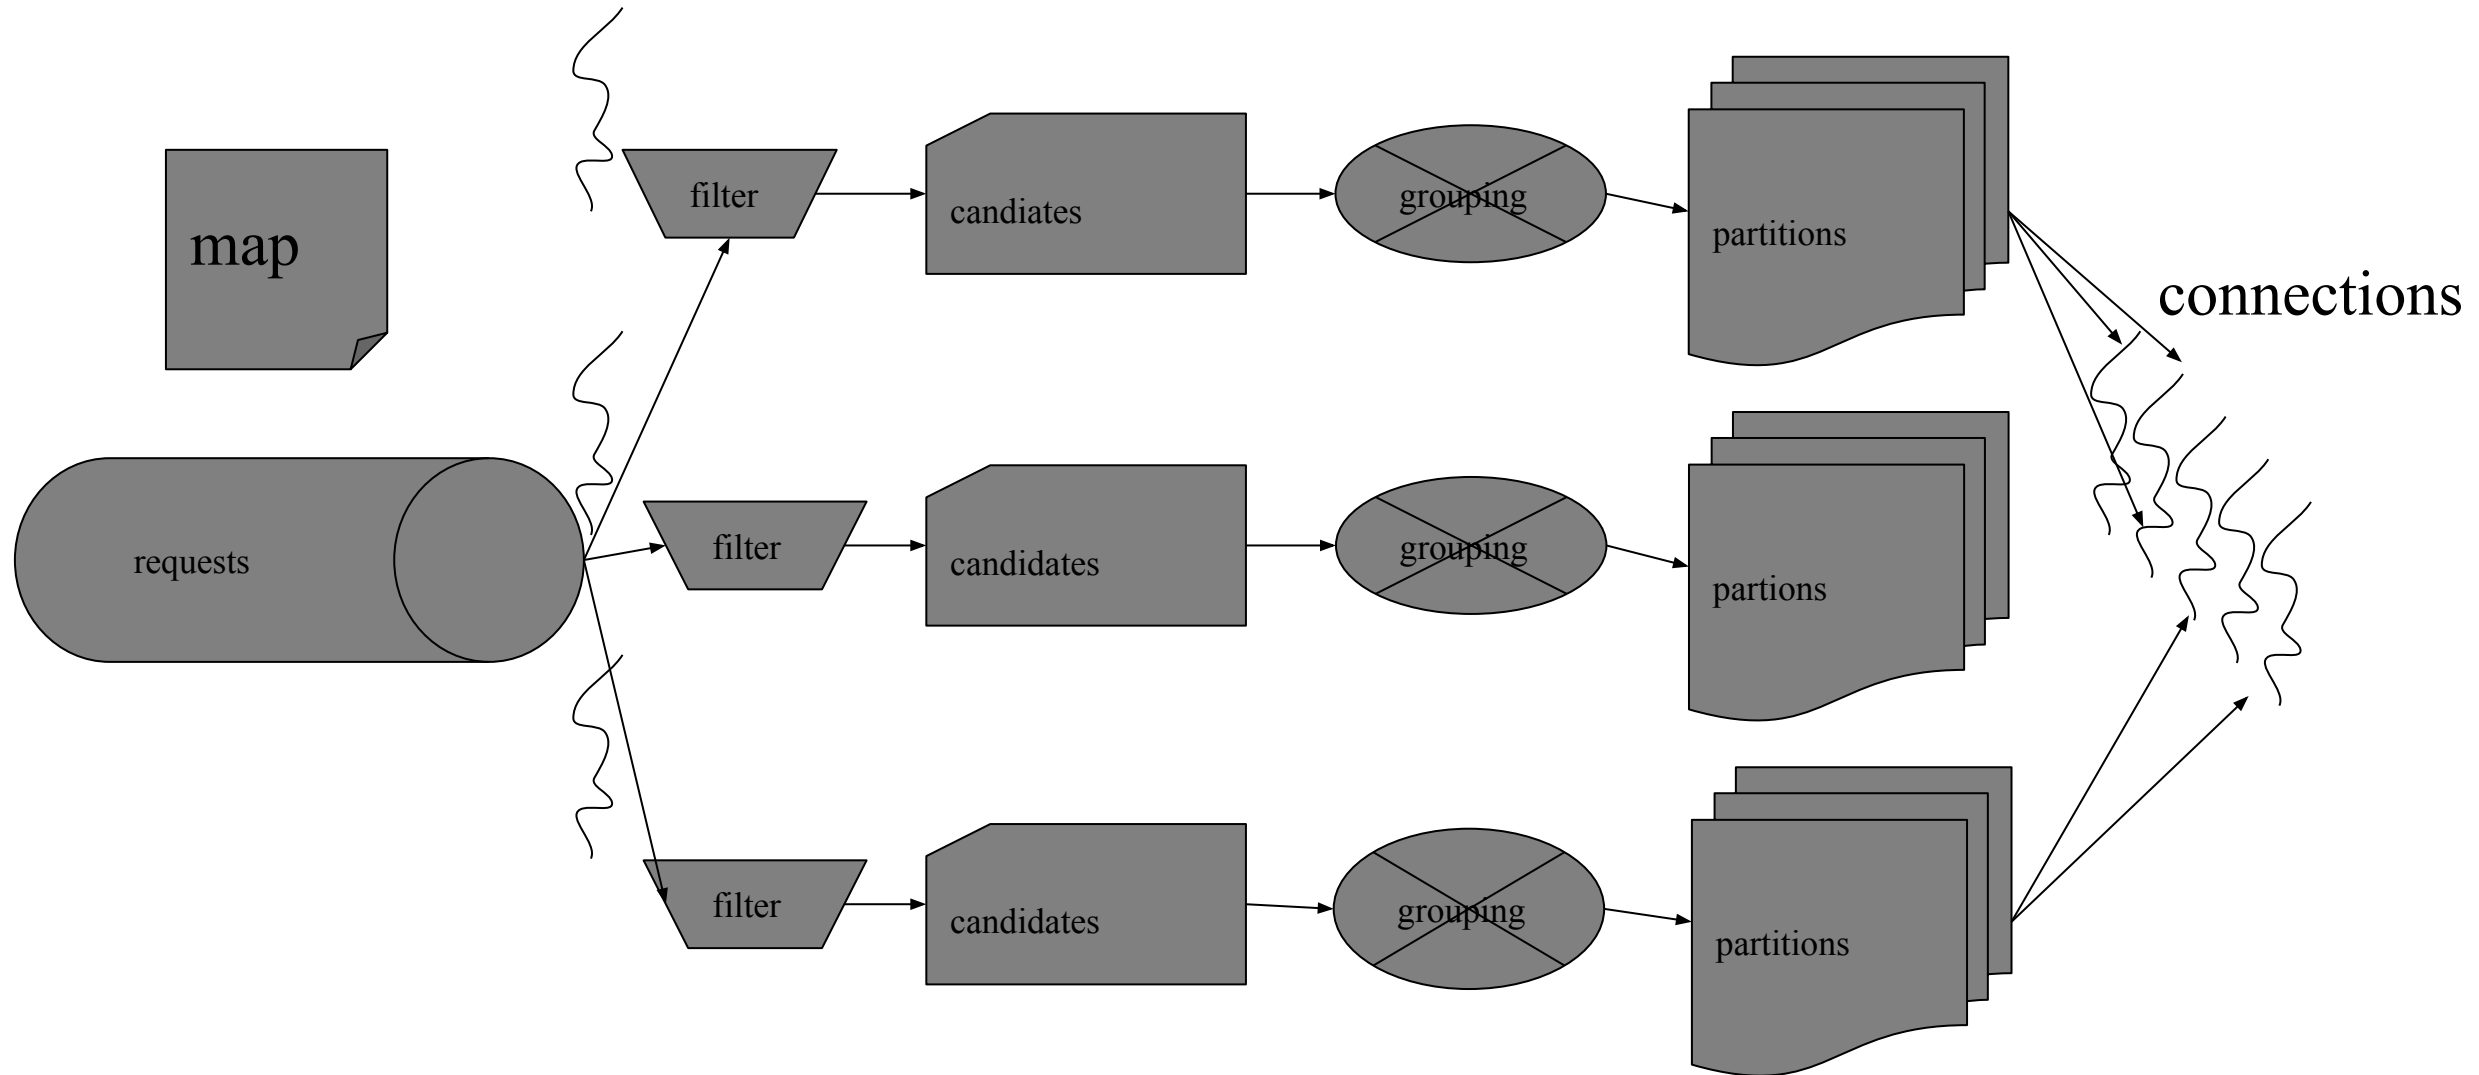
■ If $n = 2m \log_2 m$, then

$$f = \frac{1}{m^{2 \ln 2}}. \quad (5)$$

Bloom Filters in HBase

- Bloom filter B:
 - B一個bit vector表示一個集合的元素
 - 元素Y in B: Y可能在B裡, Y不在B裡的機率是f (false positive probability)
 - 元素Y not in B: Y一定不在B裡
- HFile裡存了若干個bloom filter用來表示該file裡所有出現過的keys
- Given a key k and a Hfile h, 我們不確定要查找的k是否在h裡, HBase region server會先檢查該h的bloom filter, 若檢查後發現h沒有k, 則h就不會是要查找k所對應value的對象
- Cached在region servers的記憶體裡


Client-Side Connections and Framework



HBase的適用場合

- 勿將之當作traditional relational database (RDB)來使用, 在big data大海裡撈針

- Against RDB



- 放棄transactions操作, due to CAP theorem and PAXOS
- 適合連續row keys且大量的資料操作應用, 例如AI model training時的data access (適合大量且批次的讀寫)
- 極適合用來存time-series data或structured log (e.g., CSV)

Experiences and Lessons

- HBase is a good solution conquering Moore's Law
- HBase is “extremely” flexible, but complex: dozens of parameters/design options to consider
- To manipulate a distributed database management system in a production state (specifically for big data) needs a team
 - Our lab: HBase storage engine (4 persons), Phoenix SQL (3) and R (2) each yr
- **Issues:** speed vs space, bug fix, extra functionalities expansion, backup and recovery, learning curve (for administrators and application developers), ...
- **Cooperation:** ITRI, III, Delta, ETC, UMC and ASE

Thank You

Advanced (1/5): Load Balancer

- Makes decisions for placement and movement of regions across RegionServers
 - Cluster-wide load balancing will occur only when there are no regions in transition and according to a fixed period of a time
- By default, being executed every five minutes
 - `hbase.balancer.period`
- A number of load balancers implemented:
 - configured in `hbase.master.loadbalancer.class`
 - `SimpleLoadBalancer` class
 - `FavoredNodeLoadBalancer` class
 - `StochasticLoadBalancer` class (default)

Load Balancers (cont'd)

Class Name	Description
FavoredNodeLoadBalancer	<ul style="list-style-type: none">• Assign favored nodes for each region• Roles: primary RegionServer, secondary and tertiary RegionServers
SimpleLoadBalancer	<ul style="list-style-type: none">• Invariant: number of regions each server manages shall be $\leq \text{average} \pm 1$• Given a cost function $F(C) \Rightarrow x$, the cluster will randomly try and mutate to C_{prime}• If $F(C_{\text{prime}}) < F(C)$, then switch the cluster to C_{prime}• Cost function $F()$ refers to:<ul style="list-style-type: none">• region load• table load• data locality• memstore sizes• storefile sizes
StochasticLoadBalancer	

Split Policies (cont'd)

Class Name	Description
<code>ConstantSizeRegionSplitPolicy</code>	<ul style="list-style-type: none">• Split a region as soon as any of its store files exceeds a maximum configurable size
<code>IncreasingToUpperBoundRegionSplitPolicy</code>	<ul style="list-style-type: none">• Split size is increased proportionally to (number of regions in a column store)³
<code>DelimitedKeyPrefixRegionSplitPolicy</code>	<ul style="list-style-type: none">• Group rows by a prefix of the row-key with a delimiter
<code>KeyPrefixRegionSplitPolicy</code>	<ul style="list-style-type: none">• Group rows by a prefix of the row-key
<code>DisabledRegionSplitPolicy</code>	<ul style="list-style-type: none">• Disables region splits

HMaster

- Responsible for administrative commands
- Load balancing by migration regions
- Handle failures of region servers and regions recovery
- Note: HMaster checks the aliveness of region servers through Zookeeper
 - Region servers heartbeat w

